

Levels of Measurement in statistics: -

- Nominal - Order Ordinal } \rightarrow Discrete

- Interval - Ratio } \rightarrow Continuous

Input image $\rightarrow f(i, j)$ Output image $\rightarrow g(i, j)$
Kernel $\rightarrow h(i, j)$

If $h(i, j)$ is an averaging kernel with some K .

$$g(i, j) = \frac{1}{(2K+1)^2} \sum_{u=-K}^K \sum_{v=-K}^K f(i+u, j+v)$$

\rightarrow uniform weights

a kernel with some K has size of $(2K+1) \times (2K+1)$.

for a general filter.

$$g(i, j) = \sum_{u=-K}^K \sum_{v=-K}^K h(u, v) \times f(i+u, j+v)$$

\rightarrow non-uniform weights

\downarrow
Called cross correlation.

A lot of information is lost, mainly of the edge of the image.

\rightarrow this is due to the fact that we have to get the edge pixels of the output image, hence we increase the input image by 1 row + 1 col. to obtain it. This was the case for $K=1$.
For larger K we would lose more information from the edge.

Median Filtering: - Instead of averaging we use median

- reduces noise and preserves edge.
- there are no kernels; just filter width.

Gauss - Correlation :-

$$G(i, j) = \sum_{u=-K}^K \sum_{v=-K}^K h(u, v) f(i+u, j+v)$$

$$G = h \otimes f.$$

Filtering an image \rightarrow replace each pixel with L.C. of its neighbourhood pixels. Used to do with a kernel/mask.

Box filter : \Rightarrow ① Constant filters pixel values.

• Blurry images / edges.

②

Gaussian filter : - • Bright in the middle and dark towards the edges. like the gaussian distribution in the heat map.

or Normal
 \downarrow

• Detailed edges.

parameter $\rightarrow \sigma \rightarrow$ tells us the 'pointiness' of the heat map.
variance of kernel.

more σ ; more blurry.

③ Applying a kernel on an impulse image, the result we get is basically a 2 times flipped version of the kernel in the image.

Convolution :-

$$G = h * f.$$

Gross correlation vs convolution.

\oplus

\ominus



h =

a	b	c
d	e	f
g	h	i

-1, -1	0, 1	1, -1
-1, 0	0, 0	1, 0
-1, 1	0, 1	1, 1

K=1



Convolution :- $F * G = G * F$
 $(F * G) * H = F * (G * H)$

$F = [\dots 0, 0, 1, 0, 0 \dots]$

$f * F = f$

Separable \rightarrow b/w rows & columns.

Linear Filters :-


Sharpening: \rightarrow ~~convolut~~ convolutions - averaging.

Edge Detection:

\rightarrow a rapid change in image intensity.

Edge operators should give us edge position, magnitude, orientation.

1D edge detection :-



local extrema in
1st derivative tells
us where the edges are. ~~the~~
Height of the peak tells us ~~the~~ the
strength of the edges.

2D images :- we use gradients.

$$\nabla I = \left(\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right)$$

Vertical edge.

$$\nabla I = \left(\frac{\partial I}{\partial x}, 0 \right)$$

Horizontal edge

$$\nabla I = \left(0, \frac{\partial I}{\partial y} \right)$$

Slant edge :-

$$\nabla I = \left(\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right)$$


magnitude of gradient is strength of the edge :-

$$S = \|\nabla I\| = \sqrt{\frac{\partial I^2}{\partial x} + \frac{\partial I^2}{\partial y}}$$

Orientation $\theta = \tan^{-1} \left(\frac{\partial I / \partial y}{\partial I / \partial x} \right)$.

Discrete Images

We basically focus on discrete images since pixel values are discrete and not continuous).



finding $\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y}$.

$$\frac{\partial I}{\partial x} \approx \frac{1}{2\epsilon} [(I_{i+1,j+1} - I_{i,j+1}) + (I_{i+1,j} - I_{i,j})]$$

similarly for $\frac{\partial I}{\partial y}$.

Kernel = $\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$

$$\frac{\partial}{\partial x} \approx \frac{1}{2\epsilon} \begin{bmatrix} -1 & 1 \\ -1 & 1 \end{bmatrix}$$

$$\frac{\partial}{\partial y} \approx \frac{1}{2\epsilon} \begin{bmatrix} 1 & 1 \\ -1 & -1 \end{bmatrix}$$

⇒ Hence now we have 3 edge operators.

Roberts

$$\frac{\partial}{\partial x} \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

Prewitt

$$\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ 1 & 0 & 1 \end{bmatrix}$$

Sobel

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$\frac{\partial}{\partial y} \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Good localization

Noise ineffective

Poor detection

Poor localization

Noise effective

Good detection

Edge threshold :-

Standard

Single (T)

edge

no edge

$$||\nabla I(x,y)|| > T$$

$$||\nabla I(x,y)|| < T$$

Hysteresis Based :

Two thresholds T_0 & T_1 .


$$T_0 < T_1$$

$$||\nabla I(x,y)|| < T_0 \quad \text{No edge}$$

$$||\nabla I(x,y)|| > T_1 \quad \text{Def. edge}$$

$T_0 < I < T_1$ is an edge if neighbouring pixel is an edge.

Edge Detection Using Laplacian :-



$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

Orientation is not given by L.O.

$$\frac{\partial^2 I}{\partial x^2} \approx \frac{1}{\epsilon^2} (I_{i-1,j} - 2I_{i,j} + I_{i+1,j})$$

$I_{i-1,j}$		
$I_{i,j}$	$I_{i,j}$	
		$I_{i+1,j}$

$$\frac{\partial^2 I}{\partial y^2} \approx \frac{1}{\epsilon^2} (I_{i,j-1} - 2I_{i,j} + I_{i,j+1})$$

How will this work well with continuous signals?

or $\nabla^2 \approx \frac{1}{6\epsilon^2}$

How did this modification help us?

0	1	0
1	-4	1
0	1	0

1	4	1
4	-20	4
1	4	1

For reducing noise in the image for edge detection .
 we first do $\nabla (n_r)$ Gaussian convolution of width then apply the filter to the image .
 Similarly we first do $\nabla^2 (n_r)$ then apply to the filter .
 Laplacian of the Gaussian .

Canny Edge Detector : -

1. Smooth Image with convoluting with Gaussian .
2. Take out ∇ of the smoothed out image .

Binary Images

Dilation : - Expanding the 1's to increase the area.

↳ Filling holes and gaps.

Erosion : - Shrinking the shape. creating boundaries and holes.

opening : - erosion then by dilation.
↳ is idempotent. use a disc or a square.

closing : - dilation followed by erosion.

Input : - Camera, Lidar, → Images.

Output : - Cone position in car frame, Confidence, Class

→ ~~Big~~ ~~Small~~

Big/Small
↓
colour orange blue/yellow

Lidar → gives depth of every pixel.

img = $[[0, 1, 0], [0, 1, 0], [0, 1, 0]]$

0	1	0
0	1	0
0	1	0

= $[[0, 0, 0], [1, 0, 0], [1, 0, 0]]$




(x,y) in an image is y,x in array.

$\left[\begin{array}{l} \overset{255, 255}{(255, 0, 0)}, (0, 0, 0), (0, 0, 0) \\ \overset{255, 255}{(255, 0, 0)}, (0, 0, 0), (0, 0, 0) \\ \overset{255, 255}{(255, 0, 0)}, (0, 0, 0), (0, 0, 0) \end{array} \right]$

0	0	0
1	0	0
1	0	0

(x, y, w, h)

$$\frac{\sqrt{w^2 + h^2}}{2}$$



~~Wait~~ $m = \tan \theta = \frac{h}{w}$

$$c = b + \frac{ah}{w}$$

$$y = -\frac{h}{w}x + c$$

$$\left(a - \frac{w}{2}, b + \frac{h}{2} \right)$$

$$\left(a + \frac{w}{2}, b - \frac{h}{2} \right)$$

$$\begin{aligned} & [0, 1, 0] - [0, 1, 1] \\ & ([0, 1, 0] - [0, 1, 2]) / 2, ([0, 1, 1] + [0, 1, 3]) / 2 \end{aligned}$$

Zero Padding:-

- $n \times n$ image
- $f \times f$ filter

output image $\rightarrow (n-f+1) \times (n-f+1)$

- adding zeros to the ~~output~~ input image.

- In the above eg. we have to add $f-1$ rows/columns to the input image.

until
victory
always

UNTIL
VICTORY
ALWAYS

UNTIL
V
A
W

UNL
VIC
ALWAYS

Bias:- basically a threshold for each neuron.
Bias determines the flexibility of the model.

$$g (w_1 x_1 + \dots + w_n x_n + b)$$

\downarrow \downarrow \downarrow
 activation weighted bias
 function sum

Learnable Parameters:-

Monocular camera gives \rightarrow left image.

Pointcloud \rightarrow a set of points in space made by a 3D scan.
each \rightarrow gives us a set of an image which can be moved in 3D.

Each point contains information like (x, y, z) of that point in the image frame, orientation of the point, and sometimes time of the time when the scan was done, and sometimes the RGB value of the point (pixel) \rightarrow Intensity too.

Stereo Pipeline

Right \leftrightarrow Left image \leftrightarrow Right image
 \downarrow
 disparity \rightarrow distance b/w same objects in the 2 images.

$$\text{Depth} = \frac{\text{focal length} \times \text{baseline}}{\text{disparity}} \rightarrow \text{dist. b/w left and right cameras.}$$



SIFT-?

Batch-normalization! - Normalization.
↳ scaling
the data set down

Standardization
↳ $Z_i = \frac{x_i - \bar{x}}{\sigma}$

Generalizing

Supervised

Unset Unsupervised

Deep

How to use semi-supervised learning :-

using a Blue cone we can change the RGB values

1500
Blue

300
Yellow.

Weighted loss function-

freezing a bunch of layers and then modifying other layers
to

OpenCV :-

Learnable parameters :- weights and biases.

Bias \rightarrow Each neuron has a bias.

\hookrightarrow sort of a threshold.



used to shift

the threshold \rightarrow see eg. of relu function.

$$g(w_1x_1 + \dots + w_nx_n + \beta)$$

\downarrow
bias.

relu



\rightarrow to narrow down values to or
congregate them.

• in ANN.

L.P. = $\underbrace{\text{input} \times \text{output} + \text{biases}}_{\substack{\text{no. of nodes} \\ \text{amount of weights}}} \Rightarrow \text{no. of nodes.}$

\downarrow
no. of nodes

• in CNN.

$\underbrace{\text{input}}_{\substack{\text{dense} \\ \text{no. of nodes}}} \times \underbrace{\text{output}}_{\substack{\text{filters} \times \text{size}}} + \underbrace{\text{biases}}_{\substack{\text{no. of filters}}}$

\downarrow
no. of filters

Regularization : \rightarrow reduce overfitting. training \downarrow generalization \uparrow .

L2


Batch-Size : \rightarrow no. of samples $\&$ passed through network at a time.

total = 1000 b.s. = 10 \Rightarrow 1 epoch = 100 batches.


b.s. \uparrow faster training \uparrow

Fine-tuning : \rightarrow Transfer learning applications.
 \rightarrow trained on 1 class. fine tune it to apply it to class 2 if they are similar.

Perception



Stereo : →



depth → disparity b/w L.I. and R.I.

colour → R & B values

L.I. R.I.

~~Imposing them on one another.~~

bounding boxes & key-points on L.I. using YOLO v5

Key points on R.I. using propagation → patch regression

using keypoints to draw bounding boxes on R.I.

~~calculating disparity~~ correcting them using SIFT.

Yolo v5 : Bounding box

Left image

↓ YOLO v5

BB

↓ RektNet

Keypoints

Perspective Transform

Patch Regression

RektNet : (x,y) of keypoints on every conc.

Can not run YOLO v5 on both images because the data given by YOLO is random. There is no data correlation.



sift
finds key point on both images and correlates the key-points.
then finds disparity between alone key point and SIFT key points

Problems \rightarrow time.

Mono \rightarrow faster (only 1 image) low accuracy.

PnP \rightarrow Mono using BB height.

Homogeneous
Transformation
Matrix.



$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} u \\ v \\ w \end{bmatrix} \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x_g \\ y_g \\ z_g \end{bmatrix}$$

\downarrow \downarrow \downarrow
 coordinates in Cam. Image coord in
 frame int. Matrix global frame